



Recitation 8

Devaditya Mohanty

DS-GA 1011, Fall'25
October 29, 2025

Agenda

- Efficiency Challenge
- Quantization
 - What is quantization and how to quantize?
- Pruning
 - Pruning before, during and after training?
- Knowledge Distillation
 - Distillation on outputs, weights and features


Efficiency Challenge

- What is inference? __
- Size of the model makes __



What is quantization?

- In general, it is the process of mapping a continuous or large set of values to a smaller, discrete set of values
- It is simplifying data by **rounding it to low-precision data types** while **preserving important details**



Adapted from [2]

How does quantization look like in real-world?

Original Image



“Quantized” Image




Using only 8 colors to represent the image. Adapted from [2].
Reducing the image to use just 8 colors leads to a **loss of detail and precision**

How to represent numerical values?

- Represented by “**bits**”, or binary digits
- Bits use **sign, exponent and fraction (mantissa)** to represent a value


Float 16-bit (FP16)



Adapted from [2]

What is precision? Common data types.. ?

- **Precision** is a measure of how precisely a number can be represented
- What is the difference between BF16 and FP16?



Adapted from [2]

Common data types

$(-1)^{\text{sign}} \times \text{base}^{\text{exponent}} \times \text{fraction}$

Float 32-bit (FP32)

 0 100000000 1001001000011111011011

$$(-1)^0 \times 2^1 \times 1.5707964 = 3.1415927410125732$$

higher precision

Float 16-bit (FP16)

 0 10000 1001001000

$$(-1)^0 \times 2^1 \times 1.5703125 = 3.140625$$

lower precision


original value

3.1415927

Common data types

FP32 Sign (1 bit) Exponent (8 bits) Significand / Mantissa (23 bits)



0 10000000 10010010000111111011011011





FP16 Sign (1 bit) Exponent (5 bits) Significand / Mantissa (10 bits)

0 10000 1001001000

Common data types



Common data types



(signed) INT8

(1 bit)	(7 bits)
0	1001000

Can convert FP32 to 8-bit integer representations to save a lot of memory. From [2]

How much memory do we use.. ?

- **Memory = [# bits/ 8] * # params**
- Let's say you have a *Llama-3-8B*, how much memory would you need to just load the model weights?
 - **FP32** -> 32-bits = $32/8 * 8B \sim 32$ GB
 - **FP16** -> 16-bits = $16/8 * 8B \sim 16$ GB
 - **BF16** -> 16-bits = $16/8 * 8B \sim 16$ GB
 - **INT-8** -> 8-bits = $8/8 * 8B \sim 8$ GB
 - **4-bit** -> 4-bits = $4/8 * 8B \sim 4$ GB



*Actual GPU memory usage when running a model depends on several other factors such as: CUDA/runtime overhead, activation buffers, attention caches (for generation), context size etc, and if training — gradients and optimizer/momentum states (can multiply memory by 2-6x)

Absolute maximum (absmax) quantization


- How would you scale a number into the range [-127, 127] ?
 - Divide by its absolute maximum value and multiply by 127

$$x_{\text{dequant}} = \frac{\max |X|}{127} \cdot x_{\text{quant}}$$

$$x_{\text{quant}} = \text{round}\left(\frac{x}{\max |X|} \cdot 127 \right)$$


Quantization error

- Applying the quantization and then dequantization process to retrieve the original – leads to quantization error (= delta b/w original and dequantized)



GPTQ

- Each layer is quantized independently. Given a *layer l* with a weight matrix, W_l find the quantized weights \widehat{W}_l [3]




$$\widehat{W}_l^* = \arg \min_{W_l} \|W_l X - \widehat{W}_l X\|^2$$

Model Pruning

- **What is model pruning?**
 - Removing parts of a neural network that contribute little to its output
 - Reduces model size and speeds up inference – without affecting output quality much
- **How is it different from quantization?**
 - Pruning reduces model size by reducing # parameters, while quantization does so by reducing precision of each parameters

Pruning before training

- **Procedure:** Prune the randomly initialized network followed by training from scratch. Reduces pre-training cost significantly
- **Implementation:** Weights are zeroed out using a binary mask




Adapted from [8]

- **Examples:**
 - **SNIP** [4] – one-shot pruning before training based on sensitivity
 - **GraSP** [5] – gradient-based pruning preserving training signal flow
 - **RANDOM** [6] – random pruning performs strikingly well

Pruning during training


- **Procedure:** Jointly prune and train the weights and the mask during training.



- **Examples:**
 - **Sparse Regularization** - To impose a sparse constraint on the loss function, for example adding a L_1 regularization term [7]
 - **Score based methods** - Evaluate the scores of parameters during model training and directly prune parameters with low scores [8]
 - **Sheared Llama** [12] - performs structured pruning interleaved with training

Pruning after training

- **Procedure:** Jointly prune and train the weights and the mask during training
 - Follows a pretrain-prune-retrain process, as shown in the figure below.




* Dashed lines indicate that subsequent operations are not required

- **Examples:**
 - **Lottery Ticket Hypothesis [9]:** Every overparameterized network has a smaller winning subnet that can perform just as well when trained independently
 - **Sparse GPT [10]:** LLMs can be pruned to 50% sparsity in one-shot, without any retraining
 - **Early BERT [11]:** Identifies structured winning tickets in the early stage of BERT

Knowledge Distillation


- Goal: To **transfer knowledge** from a **larger model** to a **student model**



- Why not train a student model **from scratch** instead?


Knowledge Distillation

- Simplest way is to **match the outputs** using a **distance metric**




Knowledge Distillation

- Match the **weights in intermediate layers** using a **distance metric**



Knowledge Distillation

- Match the **intermediate feature maps** using a distance metric. Ex. DistilBERT




Knowledge Distillation in LLM Era

- **Inaccessible parameters**, but we still want to transfer knowledge from LLMs




Knowledge Distillation in LLM Era

- Steer the LLM for a target skill or domain




Knowledge Distillation in LLM Era

- Feed the LLM with a small dataset as seed knowledge





Knowledge Distillation in LLM Era

- Generate knowledge from the LLM with a small dataset as seed knowledge



Knowledge Distillation in LLM Era

- Can also distill on reasoning steps (rationales) extracted from LLMs



Key Takeaways

- **Efficient inference** is important to **deploy models for resource-constrained devices**
- **Quantization** aims to simplify data by **rounding it to low-precision data types** while preserving information
- **Pruning** aims to **remove redundancy** in networks before, during or after training
- **Knowledge distillation** distills knowledge from a **teacher to student network**

References

1. Adapted from NYU DSGA-1011 (Fall'24) section material by Divyam Madaan
2. <https://newsletter.maartengrootendorst.com/p/a-visual-guide-to-quantization>
3. E. Frantar, S. Ashkboos, T. Hoefler, and D. Alistarh, “GPTQ: Accurate post-training quantization for generative pre-trained transformers,” ICLR, 2023
4. **SNIP** - N. Lee, T. Ajanthan, and P. H. S. Torr, “**SNIP**: Single-shot network pruning based on connection sensitivity,” ICLR, 2019. Available: <https://arxiv.org/abs/1810.02340>
5. **GraSP** - C. Wang, G. Zhang, and R. Grosse, “Picking winning tickets before training by preserving gradient flow,” ICLR, 2020. Available: <https://arxiv.org/abs/2002.07376>
6. S. Liu, et. al., “The unreasonable effectiveness of **random** pruning: Return of the most naive baseline for sparse training,” ICLR 2022. Available: <https://arxiv.org/abs/2202.02643>
7. W. Wen, C. Wu, Y. Wang, Y. Chen, and H. Li, “Learning structured sparsity in deep neural networks,” in NIPS, 2016.
8. H. Cheng, M. Zhang, and J. Q. Shi, “A survey on deep neural network pruning: Taxonomy, comparison, analysis, and recommendations,” arXiv preprint arXiv:2308.06767, 2024. [Online]. Available: <https://arxiv.org/abs/2308.06767>
9. J. Frankle and M. Carbin, “**The lottery ticket hypothesis**: finding sparse, trainable neural networks,” in ICLR, 2019
10. E. Frantar and D. Alistarh, “**SparseGPT**: Massive language models can be accurately pruned in one-shot,” arXiv preprint arXiv:2301.00774, 2023
11. X. Chen, Y. Cheng, S. Wang, Z. Gan, Z. Wang, and J. Liu, “**EarlyBERT**: Efficient BERT training via early-bird lottery tickets,” in ACL-IJCNLP, 2021, pp. 2195–2207.
12. M. Xia, T. Gao, Z. Zeng, and D. Chen, “**Sheared LLaMA**: Accelerating language model pre-training via structured pruning,” ICLR, 2024. Available: <https://arxiv.org/abs/2310.06694>